

Verified low-level programming *embedded in F^**

Jonathan Protzenko
Nikhil Swamy

Microsoft Research
Microsoft Research

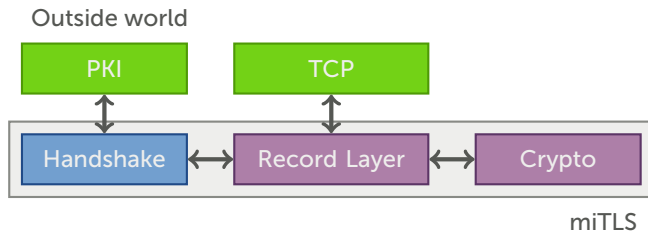
Project Everest

INRIA Paris
MSR Redmond, Cambridge, Bangalore
CMU

Part 1: Motivation

Some motivation: the TLS protocol

TLS stands for *transport layer security* and now powers more than **half** of the internet traffic.



TLS is a **cryptographic protocol**:

- "protocol": handshake, negotiation, key derivation
- "cryptographic": well, a lot of math that needs to be **fast**

Need for speed

latency is correlated with the performance of asymmetric cryptography

throughput is correlated with the performance of symmetric cryptography

Efficiency is *essential*. That's why everyone **binds** to some cryptographic library written in **C**.

Need for speed

latency is correlated with the performance of asymmetric cryptography

throughput is correlated with the performance of symmetric cryptography

Efficiency is *essential*. That's why everyone **binds** to some cryptographic library written in **C**.

Sadly, most of the time, this means, OpenSSL.

You can guess what happens next

These heavily optimized C implementations have bugs.

You can guess what happens next

OpenSSL Security Advisory [10 Nov 2016]

=====

ChaCha20/Poly1305 heap-buffer-overflow (CVE-2016-7054)

=====

Severity: High

TLS connections using *-CHACHA20-POLY1305 ciphersuites are susceptible to a DoS attack by corrupting larger payloads. This can result in an OpenSSL crash. This issue is not considered to be exploitable beyond a DoS.

have bugs.

You can guess what happens next

OpenSSL Security Advisory [10 Nov 2016]

=====

ChaCha20/Poly1305 heap-buffer-overflow (CVE-2016-7054)

=====

Severity: High

TLS connections using
attack by corrupting
issue is not consi

have bugs.

[openssl-dev] [openssl.org #4482] Wrong results with Poly1305 functions

Hanno Boeck via RT rt at openssl.org

Fri Mar 25 12:10:32 UTC 2016

- Previous message: [\[openssl-dev\] \[openssl.org #4480\] PATCH: Ubuntu 14 \(x86_64\): Compile errors and warnings when using "no-asm -ansi"](#)
- Next message: [\[openssl-dev\] \[openssl.org #4483\] Re: \[openssl.org #4482\] Wrong results with Poly1305 functions](#)
- Messages sorted by: [\[date \]](#) [\[thread \]](#) [\[subject \]](#) [\[author \]](#)

Attached is a sample code that will test various inputs for the Poly1305 functions of openssl.

These produce wrong results. The first example does so only on 32 bit, the other three also on 64 bit.

You can guess what happens next

OpenSSL Security Advisory [10 Nov 2016]

=====

ChaCha20/Poly1305 heap-buffer-overflow (CVE-2016-7054)

=====

Severity: High

TLS connections using ChaCha20/Poly1305 are vulnerable to a heap buffer overflow attack by corrupting the state of the Poly1305 function. This issue is not considered a critical severity because it only affects connections using ChaCha20/Poly1305.

have bugs.

[openssl-dev] [openssl.org #4482] Wrong results with Poly1305 functions

Hanno Boeck via RT [rt at openssl.org](#)

Fri Mar 25 12:10:32 UTC 2016

- Previous message: [\[openssl-dev\] \[openssl.org #4482\] Wrong results with Poly1305 functions](#)
- Next message: [\[openssl-dev\] \[openssl.org #4482\] Wrong results with Poly1305 functions](#)
- Messages sorted by: [\[date\]](#) [\[thread\]](#) [\[subject\]](#) [\[author\]](#)

Attached is a sample code that demonstrates the Poly1305 functions of openssl.

These produce wrong results. The other three also on 64-bit systems.

[openssl-dev] [openssl.org #4439] poly1305-x86.pl produces incorrect output

David Benjamin via RT [rt at openssl.org](#)

Thu Mar 17 21:22:26 UTC 2016

- Previous message: [\[openssl-dev\] \[openssl.org #4439\] poly1305-x86.pl produces incorrect output](#)
- Next message: [\[openssl-dev\] \[openssl.org #4439\] poly1305-x86.pl produces incorrect output](#)
- Messages sorted by: [\[date\]](#) [\[thread\]](#) [\[subject\]](#) [\[author\]](#)

Hi folks,

You know the drill. See the attached poly1305_test2.c.

```
$ OPENSSL_ia32cap=0 ./poly1305_test2
PASS
$ ./poly1305_test2
Poly1305 test failed.
got:      2637408fe03086ea73f971e3425e2820
expected: 2637408fe13086ea73f971e3425e2820
```

I believe this affects both the SSE2 and AVX2 code. It does seem to be dependent on this input pattern.

This was found because a run of our SSL tests happened to find a problematic input. I've trimmed it down to the first block where they disagree.

I'm probably going to write something to generate random inputs and stress all your other poly1305 codepaths against a reference implementation. I've seen a few other implementations, but none that I trust.

Why the sorry situation?

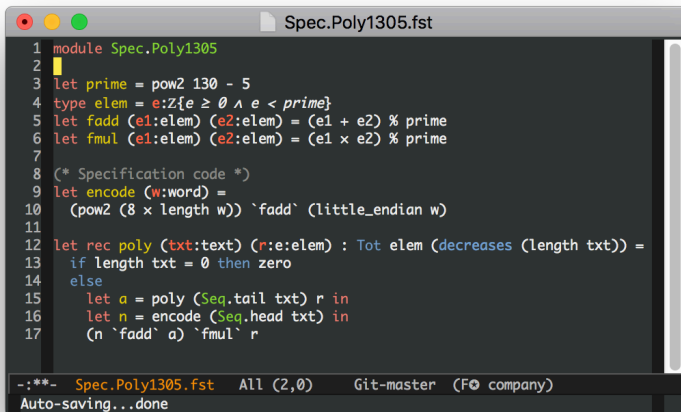
In essence: cryptography is **hard**.

- Need to understand **complex math**
- Need to optimize **complex math**
- Need to be familiar with **compilers**
- Need to reason about **side-channel resistance**

No wonder it's easy to get wrong.

Write C programs in F^* and prove them correct!

The essence of Low* in Emacs



The screenshot shows an Emacs editor window with the title bar 'Spec.Poly1305.fst'. The code is written in F* and includes comments in green. The code defines a module Spec.Poly1305, a prime number, a type elem, and functions fadd and fmul. It also includes a specification code block and a recursive function poly. The status bar at the bottom shows the file name, mode, and other information.

```
1 module Spec.Poly1305
2
3 let prime = pow2 130 - 5
4 type elem = e:Z{e ≥ 0 ∧ e < prime}
5 let fadd (e1:elem) (e2:elem) = (e1 + e2) % prime
6 let fmul (e1:elem) (e2:elem) = (e1 × e2) % prime
7
8 (* Specification code *)
9 let encode (w:word) =
10   (pow2 (8 × length w)) `fadd` (little_endian w)
11
12 let rec poly (txt:text) (r:e:elem) : Tot elem (decreases (length txt)) =
13   if length txt = 0 then zero
14   else
15     let a = poly (Seq.tail txt) r in
16     let n = encode (Seq.head txt) in
17     (n `fadd` a) `fmul` r
```

-.**~ Spec.Poly1305.fst All (2,0) Git-master (F* company)
Auto-saving...done

```

[ @"substitute" ]
val poly1305_last_pass :
acc:felem →
Stack unit
  (requires (λ h → live h acc ∧ bounds (as_seq h acc) P44 P44 P42))
  (ensures (λ h0 h1 → live h0 acc ∧ bounds (as_seq h0 acc) P44 P44 P42
    ∧ live h1 acc ∧ bounds (as_seq h1 acc) P44 P44 P42
    ∧ modifies 1 acc h0 h1
    ∧ as_seq h1 acc == HACL.Spec.Poly1305_64.poly1305_last_pass_spec_ (as_seq h0 acc)))

```

```

[ @"substitute" ]
let poly1305_last_pass_acc =
let a0 = acc.(0ul) in
let a1 = acc.(1ul) in
let a2 = acc.(2ul) in
let open HACL.Bignum.Limb in
let mask0 = gte_mask a0 HACL.Spec.Poly1305_64.p44m5 in
let mask1 = eq_mask a1 HACL.Spec.Poly1305_64.p44m1 in
let mask2 = eq_mask a2 HACL.Spec.Poly1305_64.p42m1 in
let mask = mask0 & ^ mask1 & ^ mask2 in
UInt.logand_lemma_1 (v mask0) UInt.logand_lemma_1 (v mask2);
UInt.logand_lemma_2 (v mask0) UInt.logand_lemma_2 (v mask1); UInt.logand_lemma_2 (v mask2);
UInt.logand_associative (v mask0) (v mask1) (v mask2);
cut (v mask = UInt.ones 64 ==> (v a0 ≥ pow2 44 - 5 ∧ v a1 = pow2 44 - 1 ∧ v a2 = pow2 42 - 1));
UInt.logand_lemma_1 (v HACL.Spec.Poly1305_64.p44m5); UInt.logand_lemma_1 (v HACL.Spec.Poly1305_64.p44m1);
UInt.logand_lemma_1 (v HACL.Spec.Poly1305_64.p42m1); UInt.logand_lemma_2 (v HACL.Spec.Poly1305_64.p44m5);
UInt.logand_lemma_2 (v HACL.Spec.Poly1305_64.p44m1); UInt.logand_lemma_2 (v HACL.Spec.Poly1305_64.p42m1);
let a0' = a0 & ^ (HACL.Spec.Poly1305_64.p44m5 & ^ mask) in
let a1' = a1 & ^ (HACL.Spec.Poly1305_64.p44m1 & ^ mask) in
let a2' = a2 & ^ (HACL.Spec.Poly1305_64.p42m1 & ^ mask) in
upd_3 acc a0' a1' a2'

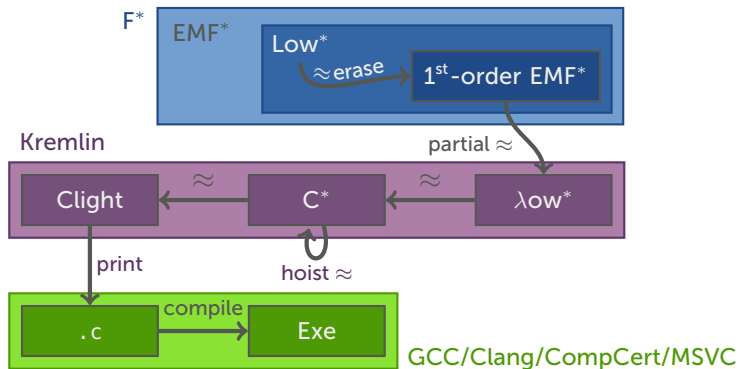
```

```

static void HACL_Impl_Poly1305_64_poly1305_last_pass(uint64_t *acc)
{
  HACL_Bignum_Fproduct carry_limb (acc);
  HACL_Bignum_Modulo_carry_top(acc);
  uint64_t a0 = acc[0];
  uint64_t a10 = acc[1];
  uint64_t a20 = acc[2];
  uint64_t a0 = a0 & (uint64_t)0xffffffff;
  uint64_t r0 = a0 >> (uint32_t)44;
  uint64_t a1 = (a10 + r0) & (uint64_t)0xffffffff;
  uint64_t r1 = (a10 + r0) >> (uint32_t)44;
  uint64_t a2 = a20 + r1;
  acc[0] = a0;
  acc[1] = a1;
  acc[2] = a2;
  HACL_Bignum_Modulo_carry_top(acc);
  uint64_t i0 = acc[0];
  uint64_t i1 = acc[1];
  uint64_t i10 = i0 & (((uint64_t)1 << (uint32_t)44) - (uint64_t)1);
  uint64_t i11 = i1 & i10 >> (uint32_t)44;
  acc[0] = i0;
  acc[1] = i1;
  uint64_t a00 = acc[0];
  uint64_t a11 = acc[1];
  uint64_t a2 = acc[2];
  uint64_t mask0 = FStar_UInt64_gte_mask(a00, (uint64_t)0xffffffff);
  uint64_t mask1 = FStar_UInt64_eq_mask(a1, (uint64_t)0xffffffff);
  uint64_t mask2 = FStar_UInt64_eq_mask(a2, (uint64_t)0x3ffffffff);
  uint64_t mask = mask0 & mask1 & mask2;
  uint64_t a0 = a00 - ((uint64_t)0xffffffff & mask);
  uint64_t a1 = a11 - ((uint64_t)0xffffffff & mask);
  uint64_t a2 = a2 - ((uint64_t)0x3ffffffff & mask);
  acc[0] = a0;
  acc[1] = a1;
  acc[2] = a2;
}

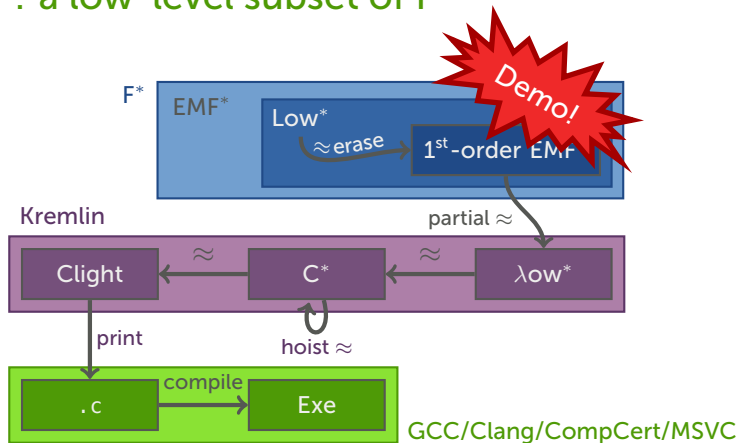
```

Low*: a low-level subset of F*



Disclaimer: these steps are supported by hand-written proofs.

Low*: a low-level subset of F*



Disclaimer: these steps are supported by hand-written proofs.

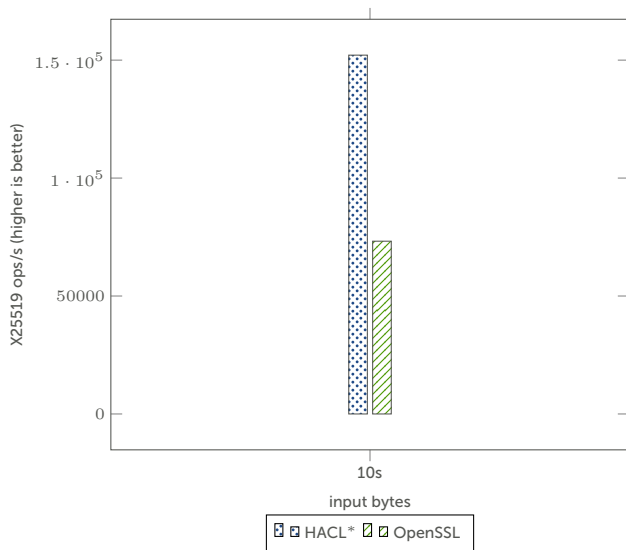
It works: HACL*

Our crypto algorithms library. Available standalone, as an OpenSSL engine, or via the NaCl API.

- Implements Chacha20, Salsa20, Curve25519, X25519, Poly1305, SHA-2, HMAC
- 7000 lines of C code
- 23,000 lines of F* code
- Performance is comparable to existing C code (not ASM)
- Some bits are in the Firefox web browser!



Jean-Karim Zinzindohoué, Karthikeyan Bhargavan,
Jonathan Protzenko, Benjamin Beurdouche
HACL*: A Verified Modern Cryptographic Library
CCS'17



Part 2: a curated subset of C

We don't need (want) all the power of C

C: hard to reason about. Instead: a *curated* subset.

Let's start with *machine integers*.

Modeling fixed-width integers

The size of `int`, `long`, `long long` types is up to the **compiler**.
Two prevalent 64-bit data models:

- LLP64, for `long long` and `t*` = 64bits (Windows)
- LP64, for `long`, `long long` and `t*` = 64bits (mostly everyone else)

That's without looking at 32-bits.

Some sanity: **C99** defines fixed-width integers like `uint32_t`, etc. in `<inttypes.h>`.

We use this revolutionary concept and expose **fixed-width integer types in Low***. (Think: curated.)

A public service announcement about undefined behavior

```
void f(long *x) {  
    long y = x = 1;  
    if (y < x) {  
        // Integer overflow, abort  
        return -1;  
    }  
    // ...  
    return 0;  
}
```

Now: `gcc -O3 -fomit-frame-pointer -S:`

```
_f:                                     ## @f  
        .cfi_startproc  
## BB#0:  
        xorl    %eax, %eax  
        retq  
        .cfi_endproc
```

A challenge of appropriately modeling C

Even when *curating* the set of features, we need to make sure the model is *restrictive* enough.

From the C11 standard, §4.3.4:

EXAMPLE An example of undefined behavior is the behavior on integer overflow.

In Low*, this means no wraparound addition for signed integer types.

For more context: “Undefined Behavior in 2017”

<https://blog.regehr.org/archives/1520>

Modeling allocations

We want **explicit control** over allocations. This means:

temporal safety an allocation is either in a stack frame, or on the heap (well-parenthesized vs. manually-managed)

spatial safety in-bounds array accesses, pointer arithmetic, no runtime length

And of course, no GC.

Another public service announcement about undefined behavior

```
int f(int *x) {  
    // C89-style  
    int foo, bar = 0;  
    int *buf = x + 16;  
  
    if (x == NULL) {  
        return -1;  
    }  
    // Proceed assuming x is non-null  
    ...  
}
```

Any decent C compiler will remove the NULL check. Ask the Linux developers.

From the C standard, §6.2.4:

If an object is referred to outside of its lifetime, the behavior is undefined.

Modeling C data layout

We offer a **predictible** compilation scheme for inductives and structures.

```
type point = {  
  x: Int32.t;  
  y: Int32.t;  
  z: Int32.t  
}
```

```
typedef {  
  int32_t x;  
  int32_t y;  
  int32_t z;  
} point;
```

- Pass by value, pay the **runtime price**
- Pass by address, pay the **verification price**

Things we don't want from C

- Addresses everywhere: a value is a value, no unfettered & operator
- Unstructured control flow (Duff's device, goto)
- Preprocessor (use meta-programming instead)
- Concurrency (maybe one day)

Things we have but I won't talk about

Many more features in support of **programmer productivity**.

- compilation of inductives to **tagged unions**
- compilation of pattern-matches to **if-then-else**
- Whole-program monomorphization
- A loop combinator library

Part 3: defining Low^*

An overview of Low* in three parts

- ① A word on C99 integers
- ② A C memory model and effects for Low*
- ③ A theory of C arrays

① A word on C99 integers

Modeling a C concept, concretely

```
module FStar.Int

  (* Generic bounded integer model, based on int *)
  let n = 32
  let max_int (n:pos) : Tot int =
    pow2 (n-1) - 1
  let min_int (n:pos) : Tot int =
    - (pow2 (n-1))
  let fits (x:int) (n:pos): Tot bool =
    min_int n <= x && x <= max_int n
  let size (x:int) (n:pos): Tot Type0 =
    b2t(fits x n)

  (* The type of integers on n bits *)
  type int_t (n:pos) = x:int{size x n}
```

Modeling a C concept, concretely (2)

```
(* Looking at the .fst file *)  
module FStar.Int32  
  
(* Instantiate the model *)  
let n = 32  
type t = | Mk: v:int_t n -> t  
  
(* For proofs only *)  
let v (x:t): GTot (int_t n) = x.v  
  
(* Expose operations *)  
let add (a:t) (b:t) : Pure t  
  (requires (size (v a + v b) n))  
  (ensures (fun c -> v a + v b = v c))  
  = Mk (add (v a) (v b))
```


Modeling a C concept, concretely (3)

```
(* Looking at the .fsti file *)  
module FStar.Int32  
  
(* Abstract type *)  
type t  
  
val v (x:t) -> GTot (int_t n)  
  
(* Lemmas re. v bijection *)  
  
(* Expose operations *)  
val add: (a:t) -> (b:t) -> Pure t  
  (requires (size (v a + v b) n))  
  (ensures (fun c -> v a + v b = v c))
```

Modeling a C concept, concretely (4)

Note that the type is *abstract*. This is essential.

- At compilation-time, we swap the model for native C integers, so we need `Int32.t </: nat`
- For soundness, the client must **not** be able to reason beyond the interface.

Note the shape of the post-condition of `add`: this is what guarantees proper overflow checking!

Fixed-width integers modules in Low*

Fixed-width machine integers are modeled in `FStar.{U,}Int{8,16,32,64,128}.fst`.

F* has integer literal syntax, in the style of C, e.g. `16ul` or `32UL`. Operators are suffixed with `^` to distinguish them from the operations on `int`.

```
module U32 = FStar.UInt32

let x = U32.( 0ul +%^ 1ul )

(* No unary minus! *)
let y = U32.( 0ul -%^ 1ul )
```

We do not have overloading (in the works).

② C memory model

Building on yesterday's discussion

We saw a **heap memory model**: a **map** from **addresses** to **values**. **Reflects** and **specifies** stateful operations in terms of pure counterparts.

```
(* Pseudo-code. *)  
val (!): #a -> r:ref a -> ... a  
  (requires (fun h -> True))  
  (ensures (fun h0 x h1 -> x = sel h0 r))
```

- The heap appears in for pre- and post-conditions
- At each program point, the **state** of the memory is reflected by a **map** from addresses to values
- The state **evolves** with program execution and so does the **map** that reflects it

Making this a C heap

We need to **restrict** yesterday's model.

yesterday automatic memory management (GC)

today manual memory management (life becomes harder)

We want to enforce **temporal safety**: no uninitialized memory accesses; no use-after-free; no double free.

Making this a C heap (2)

```
(* Reserving 0 for the NULL pointer, cf. later *)
type addr = x:nat { x > 0 }
type heap = {
  next_addr: addr;
  memory: addr -> Tot (option (a: Type0 & ... & a))
}
type ref (a:Type0): Type0 = { addr: addr; init: a; ... }

let free #a ... h r = {
  h with memory =
    fun r' -> if r' = r.addr then None else h.memory r'
}

(* Should be named: live *)
let contains #a h ... r =
  Some? (h.memory r.addr) /\
  ...
```

Making this a C heap (3)

- A live address points to **Some** `_`. A de-allocated or un-allocated address points to **None**.
- We do not allow the client to observe address reuse (important).

This heap model propagates all the way to the stateful operator, now requiring:

```
(* Pseudo-code. *)  
val (!): #a -> r:ref a -> ... a  
  (requires (fun h -> live h r))  
  (ensures (fun h0 x h1 -> x = sel h0 r))
```


Making this a C heap (3)



Demo!

- A live address points to **Some** `_`. A de-allocated or un-allocated address points to **None**.
- We do not allow the client to observe address reuse (important).

This heap model propagates all the way to the stateful operator, now requiring:

```
(* Pseudo-code. *)  
val (!): #a -> r:ref a -> ... a  
  (requires (fun h -> live h r))  
  (ensures (fun h0 x h1 -> x = sel h0 r))
```

Making this a C stack and heap

We need to go **beyond** a simple heap. We need **structure**.

About the **stack**.

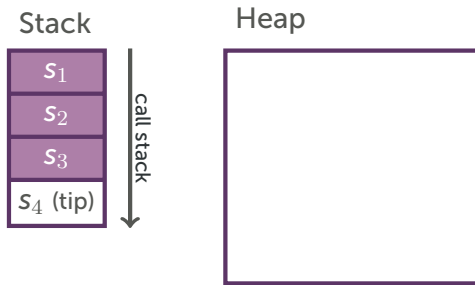
- Stack frames conceptually form a **linked list**
- Stack frames are **pushed** and **popped** on the call stack as execution progresses
- A function can only **allocate** within **its own** stack frame and should leave its calling stack **intact**
- A stack reference is only **live** as long as its stack frame is **live**

Making this a C stack and heap (2)

Stack references behave **differently**. Their **lifetime** is tied to their enclosing stack frame. They are allocated and de-allocated in **bulk**.

We are going to use **regions** to model the C stack frames, a classic framework from the 90s / 2000s.

A C memory model with regions: HyperStack



- The memory is partitioned into **stack** (s_i , left) and **heap** (the rest).
- Stack regions **nest** linearly. The topmost stack frame is the **tip**.
- Allocation is only possible in white regions.

Modeling the stack with regions (1)

The stack frames s_i are distinguished in the model. They are **transient** and may disappear. Conversely, the root and its sub-regions are **eternal**.

```
module HS = FStar.HyperStack

let root: HS.rid = HS.root

let _ =
  assert (ST.is_eternal_region root /\
    ~ (Monotonic.HyperStack.is_stack_region root))
```

Modeling the stack with regions (2)

Remember that we're still **crafting** the model.

The mechanics of the call stack are handled by **combinators** `push_frame` and `pop_frame`.

```
let push_frame (_:unit): Internal unit
  (requires (fun m -> True))
  (ensures (fun (m0:mem) _ (m1:mem) -> fresh_frame m0 m1))
  = ...

let fresh_frame (m0:mem) (m1:mem) =
  not (Map.contains m0.h m1.tip) /\
  parent m1.tip = m0.tip          /\
  m1.h == Map.upd m0.h m1.tip Heap.emp
```

The programmer does not have access to **Internal**.

Modeling the stack with regions (3)

```
let alloc_post (#a:Type) ...  
  (init:a) (m0:mem)  
  (s:mreference a ...{is_stack_region (frameOf s)}) (m1:mem)  
= is_stack_region m0.tip /\  
  Map.domain m0.h == Map.domain m1.h /\  
  m0.tip = m1.tip /\  
  frameOf s = m1.tip /\  
  HS.fresh_ref s m0 m1 /\  
  m1 == HyperStack.upd m0 s init
```

The liveness predicate is refined to capture manual lifetime in the heap, or region-based lifetime on the stack.

Modeling the stack with regions (4)

Found in `FStar.HyperStack.ST.fst`, **Stack** is the effect of well-formed client programs.

```
effect Stack (a:Type) (pre:st_pre) (post: (mem -> Tot (st_post a))) =
  STATE a (fun (p:st_post a) (h:mem) ->
    pre h /\ (forall a h1.
      (pre h /\ post h a h1 /\ equal_domains h h1) ==> p a h1))

let equal_domains (m0:mem) (m1:mem) =
  m0.tip = m1.tip
  /\ Set.equal (Map.domain m0.h) (Map.domain m1.h)
  /\ (forall r. Map.contains m0.h r ==>
    Heap.equal_dom (Map.sel m0.h r) (Map.sel m1.h r))
```

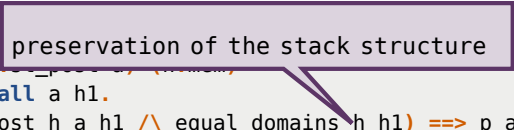
Preserves the **layout** of the stack and **doesn't allocate** in any caller frame.

Modeling the stack with regions (4)

Found in `FStar.HyperStack.ST.fst`, **Stack** is the effect of well-formed client programs.

```
effect Stack (a:T) (p:STATE a) (pre:mem) (post:mem) (post a))) =
  STATE a (fun (p:STATE a) (pre:mem) (post:mem) (post a))) =
    pre h /\ (forall a h1.
      (pre h /\ post h a h1 /\ equal_domains h h1) ==> p a h1))

let equal_domains (m0:mem) (m1:mem) =
  m0.tip = m1.tip
  /\ Set.equal (Map.domain m0.h) (Map.domain m1.h)
  /\ (forall r. Map.contains m0.h r ==>
    Heap.equal_dom (Map.sel m0.h r) (Map.sel m1.h r))
```



Preserves the **layout** of the stack and **doesn't allocate** in any caller frame.


Modeling the stack with regions (4)

Found in `FStar.HyperStack.ST.fst`, **Stack** is the effect of well-formed client programs.

```
effect Stack (a:Type) (pre:st_pre) (post: (mem -> Tot (st_post a))) =  
  STATE a (fun (p:st_post a) (h:mem) ->  
    pre h /\ (forall a h1.  
      (pre h /\ post h a h1 /\ equal_domains h h1) ==> p a h1))  
  
let equal_domains (m0:mem) (m1:mem) =  
  m0.tip = m1.tip  
  /\ Set.equal (Map.domain m0.h) (Map.domain m1.h)  
  /\ (forall r. m0.sel m0.h r ==> m1.sel m1.h r))  
  Heap. the tip remains the same
```

Preserves the **layout** of the stack and **doesn't allocate** in any caller frame.

Modeling the stack with regions (4)

Found in `FStar.HyperStack.ST.fst`,  effect of well-formed client programs.

```
effect Stack (a:Type) (pre:st_pre) (post: (mem -> Tot (st_post a))) =  
  STATE a (fun (p:st_post a) (h:mem) ->  
    pre h /\ (forall a h1.  
      (pre h /\ post h a h1 /\ equal_domains h h1) ==> p a h1))  
  
let equal_domains (m0:mem) (m1:mem) =  
  m0.tip = m1.tip  
  /\ Set.equal (Map.domain m0.h) (Map.domain m1.h)  
  /\ (forall r. Map.contains m0.h r ==>  
    Heap.equal_dom (Map.sel m0.h r) (Map.sel m1.h r))
```

Preserves the **layout** of the stack and **doesn't allocate** in any caller frame.

Recap

- Talking about the C **memory model** in F*.
- **Reflect** the semantics of stateful operations in the memory model
- **push_frame** and **pop_frame** control the call stack
- Well-formed stateful operations are in the **Stack** effect

③ A theory of C arrays

What do we need to capture...

...in order to be a faithful model of C arrays?

lifetime of variables, either tied to a stack frame, or manually-managed

spatial safety, guaranteeing access within bounds

functional model C arrays as a sequence

The Buffer library

Two flavors of buffers:

- **stack**-allocated buffers (think `alloca(3)`), lifetime tied to stack frame
- **heap**-allocated buffers (think `malloc`), must be `free`'d

Needs to be a faithful model of C:

- no **pointer arithmetic** if pointer is not live
- no zero-length allocations on the stack (use `NULL` instead)
- **length** cannot be computed at run-time
- only pointer **one-past** allowed

Beware of undefined behaviors!

The Buffer library



Two flavors of buffers:

- **stack**-allocated buffers (think `alloca(3)`), lifetime tied to stack frame
- **heap**-allocated buffers (think `malloc`), must be `free`'d

Needs to be a faithful model of C:

- no **pointer arithmetic** if pointer is not live
- no zero-length allocations on the stack (use `NULL` instead)
- **length** cannot be computed at run-time
- only pointer **one-past** allowed

Beware of undefined behaviors!

The Buffer library (1)

A buffer is **modeled** as a **reference** to a **sequence**.

```
noeq
type buffer (a: Type0) : Type0 =
| Null
| Buffer:
  (max_length: U32.t { U32.v max_length > 0 } ) ->
  (content: HST.reference (vec a (U32.v max_length))) ->
  (idx: U32.t) ->
  (length: U32.t { U32.v idx + U32.v length <= U32.v max_length
    buffer' a

(* In a given heap *)
let as_seq #a h b =
  match b with
  | Null -> Seq.createEmpty
  | Buffer max_len content idx len ->
    Seq.slice ...
```

The Buffer library (2)

```
let alloc_post_common
  (#a: Type) (r: HS.rid) (len: nat) (b: buffer a)
  (h0 h1: HS.mem): GTot Type0
= b 'unused_in' h0 /\ live h1 b /\
  (not (g_is_null b)) /\ frame0f b == r /\
  Map.domain h1.HS.h 'Set.equal' Map.domain h0.HS.h /\
  h1.HS.tip == h0.HS.tip /\ length b == len /\
  modifies_0 h0 h1

val alloca (#a: Type) (init: a) (len: U32.t):
HST.StackInline (buffer a)
  (requires (fun h -> U32.v len > 0))
  (ensures (fun h b h' ->
    alloc_post_common h.HS.tip (U32.v len) b h h' /\
    as_seq h' b == Seq.create (U32.v len) init
  ))
```

The Buffer library (3)

```
val malloc (#a: Type) (r: HS.rid) (init: a)
  (len: U32.t): HST.ST (buffer a)
  (requires (fun h -> HST.is_eternal_region r /\ U32.v len > 0))
  (ensures (fun h b h' ->
    alloc_post_common r (U32.v len) b h h' /\
    as_seq h' b == Seq.create (U32.v len) init /\
    freeable b
  ))

val free (#a: Type) (b: buffer a): HST.ST unit
  (requires (fun h0 -> live h0 b /\ freeable b))
  (ensures (fun h0 _ h1 ->
    (not (g_is_null b)) /\
    Map.domain h1.HS.h 'Set.equal' Map.domain h0.HS.h /\
    h1.HS.tip == h0.HS.tip /\
    modifies_1 b h0 h1 /\
    HS.live_region h1 (frameOf b)
  ))
```

The Buffer library (4)

Other bits from `LowStar.Buffer.fsti`:

```
val is_null (#a: Type) (b: buffer a):  
  HST.Stack bool ...  
val offset (#a: Type) (b: buffer a) (i: U32.t):  
  HST.Stack (buffer a) ...  
val index (#a: Type) (b: buffer a) (i: U32.t):  
  HST.Stack a ...  
val upd (#a: Type) (b: buffer a) (i: U32.t) (v: a):  
  HST.Stack unit ...
```

The Buffer library (4)

Other bits from `LowStar.BufferOps.fsti`:

```
(* e1.(e2) *)  
let op_Array_Access = ...  
  
(* e1.(e2) <- e3 *)  
let op_Array_Assignment = ...
```

You want to **open** this module in your code.

Reasoning with abstract modifies clauses

The heap model **precisely** defines what happens for allocations, de-allocations, modifications, etc.

However, this is **too precise** and we oftentimes want more **abstract reasoning**.

The modifies clauses library provides a set of **abstract**, **composable** predicates to talk about heap modification.

The Modifies library (1)

`LowStar.Modifies.fsti` defines an **abstract** type of memory locations which form a **monoid**.

```
val loc: Type0
val loc_none: loc
val loc_union (s1 s2: loc) : GTot loc
```

Buffers can be injected into locations (and so can regions and individual references).

```
val loc_buffer (#t: Type) (b: B.buffer t): GTot loc
```

The Modifies library (2)

`LowStar.Modifies.fsti` also defines an inclusion relation. The programmer reasons at the **granularity** of individual **buffers** or **regions**.

```
val loc_includes_region_buffer
  (#t: Type)
  (s: Set.set HS.rid)
  (b: B.buffer t)
: Lemma
  (requires (Set.mem (B.frameOf b) s))
  (ensures (loc_includes (loc_regions s) (loc_buffer b)))
  [SMTPat (loc_includes (loc_regions s) (loc_buffer b))]
```


The `Modifies` library (3)

Equipped with `locations` and the `includes` relation, the module defines `modifies` clauses.

```
val modifies (s: loc) (h1 h2: HS.mem): GTot Type0

val modifies_buffer_elim (#t1: Type) (b: B.buffer t1)
  (p: loc) (h h': HS.mem): Lemma
  (requires (
    loc_disjoint (loc_buffer b) p /\
    B.live h b /\ ... /\
    modifies p h h'
  ))
  (ensures (
    B.live h' b /\ (
      B.as_seq h b == B.as_seq h' b
    ))
  )
[ SMTPat ... ]
```

The **Modifies** library (4)

Combinators from **LowStar.Buffer** are tagged with appropriate **modifies** clauses.

```
val upd (#a: Type) (b: buffer a) (i: U32.t) (v: a):  
  HST.Stack unit  
  (requires (fun h -> live h b /\ U32.v i < length b))  
  (ensures (fun h _ h' ->  
    (not (g_is_null b)) /\  
    modifies_l b h h' /\  
    live h' b /\  
    as_seq h' b == Seq.upd (as_seq h b) (U32.v i) v  
  ))
```

This **equational theory** allows us to reason efficiently about stateful programs.

High-level verification for low-level code

For **code**, the programmer:

- opts in the Low* **effect** to model the C stack and heap;
- uses **low-level libraries** for arrays and structs;
- leverages **combinator libraries** to get C loops;
- meta-programs **first-order** code;
- relies on **data types** sparingly.

For **proofs and specs**, the programmer:

- can use **all of F^*** ,
- prove **memory safety, correctness, crypto games**, relying on
- **erasure** to yield a first-order program.

Motto: the code is **low-level** but the verification **is not**.

Part 4: KreMLin

A compiler from F* to *readable* C

The KreMLin facts:

- about 14,000 lines of OCaml
- carefully engineered to generate **readable** C code
- essential for **integration** into existing software.

Design:

- relies on the same Letouzey-style **erasure** from F*
- one internal AST with **several** compilation **passes**
- **abstract** C grammar + pretty-printer
- small amounts of **hand-written C code** (host functions)

So far, about 120k lines of C generated.

Tooling support: killing abstraction

| |
|--|
| Abstraction = good for verification |
| No Abstraction = good for compilation |

- At the module level (**-bundle**)
- At the function level (**inline_for_extraction**)

This triggers enough compiler optimizations to **fulfill the original promise**.

Tooling support: data types

Or: “programmer productivity”.

- Tuples, inductives (tagged unions) are supported
- Four (!) different compilation schemes
- Use at your own risk (MSVC! CompCert! x86 ABI!)
- Requires:
 - monomorphization
 - implementation in KreMLin of recursive equality predicates
 - mutual recursion; forward declarations

Tooling support: misc

- Type **abbreviations**
- C **loops** (syntactic closures for bodies)
- **Removal** of **uu**___
- Optimal **visibility**
- Removal of **unused** function and data types arguments
- Passing structures **by reference**

Learning KreMLin

Our work-in-progress is online:

<https://fstarlang.github.io/lowstar/html/>

Coming up: **90-minute** lab session on Low* programming and KreMLin.

Consider **opam** installing KreMLin.

Part 5: λow^* , a
model of Low^*

Removal of ghost code (1)

- Ghost is an F^* effect
- Used only for proofs, i.e. specifications (“contagious”)
- ABSOLUTELY does not fit in Low^*
- Removed via a logical relations argument

Removal of erased (2)

- `erased a` is a **computationally-irrelevant** value
- unlike `Ghost`, can be used within **code**
- used for the `log` of operations, say, in Chacha
- the value can be used in specifications via:
`val reveal: #a -> erased a -> GTot a`

All erased values, being irrelevant, can be compiled to `()` (ML).
We remove them via a **whole-program analysis**.

From monadic to effectful semantics (3)

Explicitly-monadic F^* (POPL'17) can be translated to a primitive state semantics.



Danel Ahman, Cătălin Hrițcu, Guido Martínez, Gordon Plotkin, Jonathan Protzenko, Aseem Rastogi, Nikhil Swamy.

Dijkstra Monads for Free
POPL'17

From the user-facing Low^* to λow^* (4)

A series of (unproven) **transformations** for programmer convenience.

- going from an **expression** language to a **statement** language
- compilation of **pattern-matching**
- **structures** by value
- etc.

These are performed by the **KreMLin** tool.

The core lambda-calculus: λow^*

$$\begin{aligned}\tau &::= \text{int} \mid \text{unit} \mid \{\vec{f} = \tau\} \mid \text{buf } \tau \mid \alpha \\ v &::= x \mid n \mid () \mid \{\vec{f} = v\} \mid (b, n, \vec{f}) \\ e &::= \text{let } x : \tau = \text{readbuf } e_1 \ e_2 \text{ in } e \mid \text{let } _ = \text{writebuf } e_1 \ e_2 \ e_3 \text{ in } e \\ &\quad \mid \text{let } x = \text{newbuf } n \ (e_1 : \tau) \text{ in } e_2 \mid \text{subbuf } e_1 \ e_2 \\ &\quad \mid \text{let } x : \tau = \text{readstruct } e_1 \text{ in } e \mid \text{let } _ = \text{writestruct } e_1 \ e_2 \text{ in } e \\ &\quad \mid \text{let } x = \text{newstruct } (e_1 : \tau) \text{ in } e_2 \mid e_1 \triangleright f \\ &\quad \mid \text{withframe } e \mid \text{pop } e \mid \text{if } e_1 \text{ then } e_2 \text{ else } e_3 \\ &\quad \mid \text{let } x : \tau = d \ e_1 \text{ in } e_2 \mid \text{let } x : \tau = e_1 \text{ in } e_2 \mid \{\vec{f} = e\} \mid e.f \mid v \\ P &::= \cdot \mid \text{let } d = \lambda y : \tau_1. e : \tau_2, P\end{aligned}$$

About λow^* :

- a **type system** without progress
- in a **simulation** with the original F^* program
- standard **substitutive** semantics.

The judgements of λow^*

Typing judgement:

$$\Gamma_P; \Sigma; \Gamma \vdash e : \tau$$

where:

- Γ_P is the set of **global** program **definitions**
- Σ is the **store typing**
- Γ is the **local context**

The judgements of λow^* (2)

Reduction semantics:

$$P \vdash (H, e) \xrightarrow{\ell} (H', e')$$

where:

- ℓ is the set of **trace events**
- H is the stack of **frames**

The judgements of λow^* (2)

Reduction semantics:

$$P \vdash (H, e) \xrightarrow{\ell} (H', e')$$

where:

- ℓ is the set of **trace events**
- H is the stack of **frames**

Let's talk about traces!

What are we protecting against

- We want to guard against some **memory** and **timing** side-channels
- Our **secret** data is at an **abstract** type
- By using **abstraction**, we can **control** what operations we allow on secret data

Abstraction to the rescue

Our module for **secret integers** exposes a handful of **audited, carefully-crafted** functions that we trust have **secret-independent** traces.

```
(* limbs only ghostly revealed as numbers *)  
val v : limb -> Ghost nat  
  
val eq_mask: x:limb -> y:limb ->  
  Tot (z:limb {if v x <> v y then v z = 0 else v z = pow2 26 - 1})
```

By **construction**, the programmer **cannot** use a `limb` for branching or array accesses.

What we show

We model **trace events** as part of our reduction.

$$\ell ::= \cdot \mid \text{read}(b, n, \vec{f}) \mid \text{write}(b, n, \vec{f}) \mid \text{brT} \mid \text{brF} \mid \ell_1, \ell_2$$

Note: this does not rule out ALL side channels!

Secret-independence: an intuition

A type-indexed relation $v_1 \equiv_{\tau} v_2$ over values:

$$n \equiv_{\text{int}} n$$

$$v_1 \equiv_a v_2$$

...

Intuition: terms are related if they only differ on sub-terms at secret types.

Main theorem: functions, when applied to related values in related stores, have related reductions and **emit the same traces**.

Note: this only goes up to CompCert Clight

Theorem (Secret independence)

Given

- ① *A program well-typed against a secret interface, Γ_s , i.e., $\Gamma_s, \Gamma_P; \Sigma; \Gamma \vdash (H, e) : \tau$, where e is not a value.*
- ② *A well-typed implementation of the Γ_s interface, $\Gamma_s; \Sigma; \cdot \vdash_{\Delta} P_s$, such that P_s is equivalent modulo secrets.*
- ③ *A pair (ρ_1, ρ_2) of well-typed (related) substitutions for Γ .*

There exists $\ell, \Sigma' \supseteq \Sigma, \Gamma', H', e'$ and a pair (ρ'_1, ρ'_2) of well-typed substitutions for Γ' , such that

- ① *$P_s, P \vdash (H, e)[\rho_1] \rightarrow_{\ell}^+ (H', e')[\rho'_1]$ if and only if, $P_s, P \vdash (H, e)[\rho_2] \rightarrow_{\ell}^+ (H', e')[\rho'_2]$, and*
- ② *$\Gamma_s, \Gamma_P; \Sigma'; \Gamma' \vdash (H', e') : \tau$*

Next step: C^* , an imperative language

This our next intermediary language.

- **statement** language
- not substitutive semantics (stack of **contexts** with holes)
- expressions are **pure**
- **deterministic**

We relate λow^* programs to C^* programs via a simulation.

A glimpse of the reduction rules

From λow^* :

$$\frac{}{P \vdash (H, \text{if } 0 \text{ then } e_1 \text{ else } e_2) \rightarrow_{\text{brF}} (H, e_2)} \text{LifF}$$

From C^* :

$$\frac{\llbracket \hat{e} \rrbracket_{(V)} = 0}{\hat{P} \vdash (S, V, \text{if } \hat{e} \text{ then } \vec{s}_1 \text{ else } \vec{s}_2; \vec{s}) \rightsquigarrow_{\text{brF}} (S, V, \vec{s}_2; \vec{s})} \text{ClfF}$$

Theorem

The C^ program \hat{P} terminates with trace ℓ and return value v , i.e., $\hat{P} \vdash ([], V, \vec{s}; \text{return } \hat{e}) \xrightarrow{\ell, *} ([], V', \text{return } v)$ if, and only if, so does the λow^* program: $P \vdash (\{\}, e[V]) \xrightarrow{\ell, *} (H', v)$; and similarly for divergence.*

Next step: C* to CompCert Clight

We encode the trace preservation using **builtins** that generate **trace events**.

Two relevant bits:

- **hoisting**, which changes the memory layout (abstract traces)
- **struct passing**, which changes the memory accesses (two passes)

Final step: Clight to assembly

Some possible approaches:

- **instrument** CompCert (Barthe *et al.*)
- Use **Vellvm** (Zdancewicz *et al.*)